

Concurrency multiple computations in overlapping time periods;
does not have to be simultaneous

Parallelism multiple computations executing simultaneously

Parallel computation occurs at different level:

- spread across computers (e.g., with MapReduce)
- multiple cores of a CPU executing different instructions (MIMD)
- multiple cores of a CPU executing same instruction (SIMD)
 - e.g. GPU rendering pixels

Both parallelism and concurrency need to deal with *synchronisation*.

Example: *Map-reduce* is a popular programming model for

- manipulating very large data sets
- on a large network of computers (local or distributed)

The *map* step filters data and distributes it to nodes

- data distributed as (*key*, *value*) pairs
- each node receives a set of pairs with common *key*(s)

Nodes then perform calculation on received data items

The *reduce* step computes the final result

- by combining outputs (calculation results) from the nodes

Also needs a way to determine when all calculations completed

1

2

Parallelism Across a an Array

GPU rendering pixels

- multiple identical processors
- each given one element of a data structure from main memory
- each performing same computation on that element (SIMD)
- results copied back to main memory data structure



But not totally independent: need to *synchronise* on completion

Parallelism Across Processes

One method for creating parallelism:

Use `posix_spawn()` to create multiple processes each of which does part of job.

- child executes concurrently with parent
- runs in its own address space
- inherits some state information from parent (e.g. open fd's)

Processes have some disadvantages

- process switching expensive
- each require a significant amount of state (RAM)
- communication between processes limited and/or slow

3

4

threads - mechanism for parallelism within process.

- threads allow simultaneous execution within process
- each threads has its own execution state
- threads within a process share address space:
 - threads share code (functions)
 - threads share global & static variables
 - threads share heap (malloc)
- but separate stack for each thread
 - local variables not shared
- threads share file descriptor
- threads share signals

*// POSIX threads widely supported in Unix-like
// and other systems (Windows). Provides functions
// to create/synchronize/destroy/... threads*

```
#include <pthread.h>
```

5

6

Create A POSIX Thread

```
int pthread_create(pthread_t *thread,
                  const pthread_attr_t *attr,
                  void *(*start_routine)(void *),
                  void *arg);
```

- creates a new thread with specified attributes (can be NULL)
- thread info stored in **thread*
- thread starts by executing *start_routine(arg)*
- returns 0 if OK, -1 otherwise and sets *errno*
- analogous to *posix_spawn()*

7

Wait for A POSIX Thread

```
int pthread_join(pthread_t thread, void **retval)
```

- wait until *thread* terminates
- *thread* return (or *pthread_exit()*) value is placed in **retval*
- if *thread* has already exited, does not wait
- if main returns or *exit* called, all threads terminated
- so typically wait for all threads
- analogous to *waitpid*

8

```
void pthread_exit(void *retval);
```

- terminate execution of thread (and free resources)
- `retval` is returned (see `pthread\join`)
- if `thread` has already exited, does not wait
- analagous to `exit`

Incremeing a global variable is not an atomic (indivisible) operation.

```
int bank_account;

void *thread(void *a) {
    // ...
    bank_account++;
    // ...
}

la    $t0, bank_account
lw    $t1, ($t0)
addi  $t1, $t1, 1
sw    $t1, ($t0)
.data
bank_account: .word 0
```

9

10

If `bank_account == 42` and two threads increment simultaneously.

<pre>la \$t0, bank_account lw \$t1, (\$t0) # \$t1 == 42 addi \$t1, \$t1, 1 # \$t1 == 43 sw \$t1, (\$t0) # bank_account == 43</pre>	<pre>la \$t0, bank_account lw \$t1, (\$t0) # \$t1 == 42 addi \$t1, \$t1, 1 # \$t1 == 43 sw \$t1, (\$t0) # bank_account == 43</pre>
--	--

One increment is lost.

Note threads don't share registers or stack (local variable).

They do share global variables.

If `bank_account == 100` and two threads change simultaneously.

<pre>la \$t0, bank_account lw \$t1, (\$t0) # \$t1 == 100 addi \$t1, \$t1, 100 # \$t1 == 200 sw \$t1, (\$t0) # bank_account == ?</pre>	<pre>la \$t0, bank_account lw \$t1, (\$t0) # \$t1 == 100 addi \$t1, \$t1, -50 # \$t1 == 50 sw \$t1, (\$t0) # bank_account == 50 or 200</pre>
---	--

```
int pthread_mutex_lock(pthread_mutex_t *mutex);
int pthread_mutex_unlock(pthread_mutex_t *mutex);
```

- only one thread can enter a *critical section*
- establishes mutual exclusion — *mutex*
- call `pthread_mutex_lock` before
- call `pthread_mutex_unlock` after
- only 1 thread can execute in

```
pthread_mutex_lock(&bank_account_lock);
andrews_bank_account += 1000000;
pthread_mutex_unlock(&bank_account_lock);
```

Semaphores are special variables which provide a more general synchronisation mechanism than mutexes.

```
#include <semaphore.h>
int sem_init(sem_t *sem, int pshared,
             unsigned int value);
int sem_post(sem_t *sem);
```

13

```
#include <semaphore.h>
sem_t sem;
sem_init(&sem, 0, n);

sem_wait(&sem);
// only n threads can be in executing
// in here simultaneously
sem_post(&sem);
```

14

```
int flock(int FileDesc, int Operation)
```

Similar to mutexes for a file.

- controls access to shared files (**note:** files not fds)
- possible operations
 - LOCK_SH ... acquire shared lock
 - LOCK_EX ... acquire exclusive lock
 - LOCK_UN ... unlock
 - LOCK_NB ... operation fails rather than blocking
- in blocking mode, `flock()` does not return until lock available
- only works correctly if all processes accessing file use locks
- return value: 0 in success, -1 on failure

15

If a process tries to acquire a *shared lock* ...

- if file not locked or other shared locks, OK
- if file has exclusive lock, blocked

If a process tries to acquire an *exclusive lock* ...

- if file is not locked, OK
- if any locks (shared or exclusive) on file, blocked

If using a non-blocking lock

- `flock()` returns 0 if lock was acquired
- `flock()` returns -1 if process would have been blocked

16

Concurrency is *complex* with many issues beyond this course:

Data races thread behaviour depends on unpredictable ordering;
can produce difficult bugs or security vulnerabilities

Deadlock threads stopped because they are wait on each other

Livelock threads running without making progress

Starvation threads never getting to run