

Week 05 Tutorial Solutions

1. If the data segment of a particular MIPS program starts at the address 0x10000020, then what addresses are the following labels associated with, and what value is stored in each 4-byte memory cell?

```
.data
a: .word 42
b: .space 4
c: .asciiz "abcde"
   .align 2
d: .byte 1, 2, 3, 4
e: .word 1, 2, 3, 4
f: .space 1
```

Answer:

Memory map showing labels, associated addresses, and contents of memory cells:

Label	Address	Contents	Contents in hex
a	0x10010020	42	0x0000002A
b	0x10010024	???	0x????????
c	0x10010028	'a' 'b' 'c' 'd'	0x61626364
	0x1001002C	'e' '\0' ? ?	0x6500????
d	0x1001030	1, 2, 3, 4	0x01020304
e	0x1001034	1	0x00000001
	0x1001038	2	0x00000002
	0x100103C	3	0x00000003
	0x1001040	4	0x00000004
f	0x1001044	?	0x????????

2. Give MIPS directives to represent the following variables:

- a. int u;
- b. int v = 42;
- c. char w;
- d. char x = 'a';
- e. double y;
- f. int z[20];

Assume that we are placing the variables in memory, at an appropriately-aligned address, and with a label which is the same as the C variable name.

Answer:

Note that .space *n* allocates an uninitialised *n*-byte region of memory, while .word *v* allocates four bytes (one *word*) of space, initialised with the value *v*.

- a. u: .space 4
- b. v: .word 42
- c. w: .space 1
- d. x: .byte 'a';
- e. y: .space 8
- f. z: .space 80 (20 * 4-byte ints)

3. Consider the following memory state:

Address	Data Definition
0x10010000	aa: .word 42
0x10010004	bb: .word 666
0x10010008	cc: .word 1
0x1001000C	.word 3
0x10010010	.word 5
0x10010014	.word 7

What address will be calculated, and what value will be loaded into register `$t0`, after each of the following statements (or pairs of statements)?

a. `la $t0, aa`

b. `lw $t0, bb`

c. `lb $t0, bb`

d. `lw $t0, aa+4`

e. `la $t1, cc`
`lw $t0, ($t1)`

f. `la $t1, cc`
`lw $t0, 8($t1)`

g. `li $t1, 8`
`lw $t0, cc($t1)`

h. `la $t1, cc`
`lw $t0, 2($t1)`

Answer:

a. `la $t0, aa`

Loads the address associated with the label `aa` into `$t0`. Since `aa` is located at `0x10010000`, the value loaded into `$t0` is `0x10010000`.

b. `lw $t0, bb`

Load the contents of the 4-byte memory cell associated with the label `bb` into the register `$t0`. Since `bb` is located at `0x10010004`, the address used for this instruction is `0x10010004`. The contents of that address is 666, and this is the value loaded into register `$t0`.

c. `lb $t0, bb`

Load the contents of the 1-byte memory cell associated with the label `bb` into the register `$t0`. The address is clear, and the same as in the previous question, `0x10010004`. What's less clear is which byte of the contents (666) will be loaded.

The first thing to do is to work out what 666 looks like as a 32-bit quantity: `0x0000029a`. We need to decide whether the byte containing `0x00` or the byte containing `0x9a` is located at `0x10010004`. The answer to this depends on whether the machine is *big-endian* or *little-endian* (i.e., what order we take the bytes from a 4-byte memory word, when we access them byte-at-a-time). On CSE, the byte containing `0x9a` is loaded.

The `lb` instruction has another interesting property: it *sign-extends* the value to 32 bits. This means that it will propagate the high-order bit in the byte all the way to the high-order bit in the 32-bit word, so `$t0` ends up with the value `0xffffffff9a`, since the high-order bit of `0x9a` is a one.

d. `lw $t0, aa+4`

This loads a 32-bit value from the memory location that is 4 bytes beyond the memory location associated with the label `aa`. Since `aa` is associated with `0x10010000`, the address is determined as `0x10010004`, and the value 666 is loaded into `$t0`.

e. `la $t1, cc`
`lw $t0, ($t1)`

The first instruction loads `cc`'s address into register `$t1` (i.e., `0x10010008`). The second instruction then uses the value in `$t1` as an address, and loads the 32-bit quantity from that address. Thus the address is `0x10010008`, and the value loaded into `$t0` is `0x00000001` (i.e. 1).

The two instructions are equivalent to `lw $t0, cc`.

f. `la $t1, cc`
`lw $t0, 8($t1)`

The first instruction loads `cc`'s address into register `$t1` (i.e., `0x10010008`). The second instruction then takes the value in `$t1`, adds 8 to it, and uses the result as the address. So, the address used by the second instruction is $8 + 0x10010008 = 0x10010010$. Thus, the value 5 is loaded into `$t0`.

g.

```
li    $t1, 8
lw    $t0, cc($t1)
```

The first instruction loads the constant 8 into register `$t1`. The second instruction then takes the address associated with `cc`, adds the contents of `$t1` to it, and uses the result as the address. So, the address used by the second instruction is $0x10010008 + 8 = 0x10010010$ (the same as in the previous question). Thus, the value 5 is loaded into `$t0`.

h.

```
la    $t1, cc
lw    $t0, 2($t1)
```

The first instruction loads `cc`'s address into register `$t1` (i.e., `0x10010008`). The second instruction then takes the value in `$t1`, adds 2 to it, and uses the result as the address. So, the address used by the second instruction is $2 + 0x10010008 = 0x1001000A$.

However: because this is a `lw` instruction, the address must be 4-byte aligned. Thus, executing this instruction will result in a memory alignment error.

4. What is a breakpoint?

When is it useful in debugging?

Answer:

Discussed in tute.

5. Translate this C program to MIPS assembler

```
#include <stdio.h>

int main(void) {
    int i;
    int numbers[10] = {0};

    i = 0;
    while (i < 10) {
        scanf("%d", &numbers[i]);
        i++;
    }
}
```

Answer:

```

# read 10 numbers into an array
# i in register $t0
main:

    li $t0, 0          # i = 0
loop0:
    bge $t0, 10, end0  # while (i < 10) {

    li $v0, 5          # scanf("%d", &numbers[i]);
    syscall            #

    mul $t1, $t0, 4     # calculate &numbers[i]
    la $t2, numbers    #
    add $t3, $t1, $t2   #
    sw $v0, ($t3)       # store entered number in array

    add $t0, $t0, 1     # i++;
    b loop0            # }
end0:

    jr $31             # return

.data

numbers:
    .word 0 0 0 0 0 0 0 0 0 0 # int numbers[10] = {0};

```

6. Translate this C program to MIPS assembler

```

#include <stdio.h>

int main(void) {
    int i;
    int numbers[10] = {0,1,2,3,4,5,6,7,8,9};

    i = 0;
    while (i < 10) {
        printf("%d\n", numbers[i]);
        i++;
    }
}

```

Answer:

```

# i in register $t0
main:
    li $t0, 0          # i = 0
loop1:
    bge $t0, 10, end1  # while (i < 10) {

    mul $t1, $t0, 4     # calculate &numbers[i]
    la $t2, numbers    #
    add $t3, $t1, $t2   #
    lw $a0, ($t3)       # Load numbers[i] into $a0
    li $v0, 1          # printf("%d", numbers[i])
    syscall

    li $a0, '\n'       # printf("%c", '\n');
    li $v0, 11
    syscall

    add $t0, $t0, 1     # i++
    b loop1            # }
end1:

    jr $31             # return

.data

numbers:
    .word 0 1 2 3 4 5 6 7 8 9 # int numbers[10] = {0 1 2 3 4 5 6 7 8 9};

```

```
.word 0 1 2 3 4 5 6 7 8 9 # int numbers[10] = {0,1,2,3,4,5,6,7,8,9};
```

7. Translate this C program to MIPS assembler

```
int main(void) {
    int i;
    int numbers[10] = {0,1,2,-3,4,-5,6,-7,8,9};

    i = 0;
    while (i < 10) {
        if (numbers[i] < 0) {
            numbers[i] += 42;
        }
        i++;
    }
}
```

Answer:

```
# i in register $t0
main:

    li $t0, 1          # i = 1
loop2:
    bge $t0, 10, end2  # while (i < 10) {

    mul $t1, $t0, 4     #
    la $t2, numbers    #
    add $t3, $t1, $t2   # $t3 = &numbers[i]
    lw $t5, ($t3)       # $t5 = numbers[i]

    bge $t5, 0, skip2   # if (numbers([i]) < 0) {
    add $t5, $t5, 42    #     numbers[i] += 42
    sw $t5, ($t3)       #
                        # }
skip2:
    add $t0, $t0, 1     # i++;
    b loop2             # }
end2:

    jr $31              # return

.data

numbers:
    .word 0 1 2 -3 4 -5 6 -7 8 9 # int numbers[10] = {0,1,2,-3,4,-5,6,-7,8,9};
```

8. Translate this C program to MIPS assembler

```
#include <stdio.h>

int main(void) {
    int i;
    int numbers[10] = {0,1,2,3,4,5,6,7,8,9};

    i = 0;
    while (i < 5) {
        int x = numbers[i];
        int y = numbers[9 - i];
        numbers[i] = y;
        numbers[9 - i] = x;
        i++;
    }
}
```

Answer:

```

# i in register $t0, x in $t4, y in $t8
main:

    # assume there is code here to assign values to the array

    li $t0, 0          # i = 0
loop2:
    bge $t0, 5, end2    # while (i < 5) {

    mul $t1, $t0, 4      #
    la $t2, numbers     #
    add $t3, $t1, $t2    # $t3 = &numbers[i]
    lw $t4, ($t3)        # x = numbers[i]

    li $t5, 9           # $t5 = 9 - i
    sub $t5, $t5, $t0    #
    mul $t6, $t5, 4      #
    add $t7, $t6, $t2    # $t7 = &numbers[9 - i]
    lw $t8, ($t7)        # y = numbers[9 - i]

    sw $t8, ($t3)        # numbers[i] = y
    sw $t4, ($t7)        # numbers[9 - i] = x

    add $t0, $t0, 1      # i++;
    b loop2              # }
end2:

    jr $31               # return

.data

numbers:
    .word 0 1 2 3 4 5 6 7 8 9 # int numbers[10] = {0,1,2,3,4,5,6,7,8,9};

```

9. The following loop determines the length of a string, a '\0'-terminated character array:

```

char *string = "....";
char *s = &string[0];
int length = 0;
while (*s != '\0') {
    length++; // increment length
    s++;      // move to next char
}

```

Write MIPS assembly to implement this loop.

Assume that the variable `string` is implemented like:

```

.data
string:
    .asciiz "...."

```

Assume that the variable `s` is implemented as register `$t0`, and variable `length` is implemented as register `$t1`. And, assume that the character `'\0'` can be represented by a value of zero.

Answer:

```
.data
string:
    .asciiz "...."

.text
la    $t0, string    # s = &string[0];
li    $t1, 0
while:
lb    $t2, ($t0)      # if (*s == 0) goto end_loop
beq   $t2, $0, end_loop

addi  $t1, $t1, 1     # Length++
addi  $t0, $t0, 1     # s++
j     while           # goto while

# $t1 contains the length of the string
```

10. Conceptually, the MIPS pseudo-instruction to load an address could be encoded as something like the following:



Since addresses in MIPS are 32-bits long, how can this instruction load an address that references the data area, such as 0x10010020?

Answer:

The la pseudo-instruction is implemented as two real MIPS instructions:

```
lui    $t1, BaseAddr    # top 16 bits of address
ori    $t1, $t1, Offset  # bottom 16 bits of address
```

... so the address is actually loaded in two sections: the top 16 bits of the register are loaded with the top 16 bits of the base address for the region being referenced; then the bottom 16 bits are loaded with an offset into the region. The assembler splits the address into its two 16-bit components, and places the top 16 bits in the lui (load upper immediate) instruction, and the bottom 16 bits in the ori (bitwise-or immediate) instruction.

For example, imagine that we want to load the address of an object located at 0x10010020 in a data region that starts at address 0x10010000. The first instruction would load 0x1001 into the top 16 bits of \$t1, giving it a value of 0x10010000. The second instruction would bitwise-OR this with the offset value i.e. 0x0020. The final address would thus be produced via

Address = 0x10010000 | 0x0020 = 0x10010020

The same technique is used for the li (load immediate) instruction, which loads a 32-bit constant value into a specified register.

11. Implement the following C code in MIPS assembly instructions, assuming that the variables x and y are defined as global variables (within the .data region of memory):

```
long x;    // assume 8 bytes
int y;     // assume 4 bytes

scanf("%d", &y);

x = (y + 2000) * (y + 3000);
```

Assume that the product might require more than 32 bits to store.

Answer:

```

.data
x: .space 8
y: .space 4
.text

li    $v0, 5
syscall
sw    $v0, y

lw    $t0, y
addi  $t0, $t0, 2000
lw    $t1, y
addi  $t1, $t1, 3000
mult  $t0, $t1      # (Hi,Lo) = $t0 * $t1
mfhi  $t0
sw    $t0, x        # top 32 bits of product
mflo  $t0
sw    $t0, x+4      # bottom 32 bits of product

```

12. Write MIPS assembly to evaluate the following C expression, leaving the result in register \$v0.

```
((x*x + y*y) - x*y) * z
```

Write one version that minimises the number of instructions, and another version that minimises the number of registers used (without using temporary memory locations).

Assume that: all variables are in labelled locations in the .data segment; the labels are the same as the C variable names; all results fit in a 32-bit register (i.e., no need to explicitly use Hi and Lo).

Answer:

A version that aims to minimise instructions (using more registers):

```

lw    $t0, x
lw    $t1, y
mul   $t2, $t0, $t0
mul   $t3, $t1, $t1
add   $t4, $t2, $t3
mul   $t5, $t0, $t1
sub   $t5, $t4, $t5
lw    $t4, z
mul   $v0, $t5, $t4

```

A version that aims to minimise registers used (using more instructions):

```

lw    $t0, x
lw    $t1, y
mul   $t0, $t0, $t0
mul   $t1, $t1, $t1
add   $t0, $t0, $t1
lw    $t1, x
lw    $t2, y
mul   $t1, $t1, $t2
sub   $t0, $t0, $t1
lw    $t1, z
mul   $v0, $t0, $t1

```

An even better solution from John Luo (minimising instructions and registers):

```

lw    $t0, x      # t0 = x
lw    $v0, y      # v0 = y
mul   $t1, $t0, $t0 # t1 = x * x
mul   $t0, $t0, $v0 # t0 = x * y
mul   $v0, $v0, $v0 # v0 = y * y
add   $v0, $t1, $v0 # v0 = (x * x) + (y * y)
sub   $v0, $v0, $t0 # v0 = ((x * x) + (y * y)) - (x * y)
lw    $t0, z      # t0 = z
mul   $v0, $v0, $t0 # v0 = (((x * x) + (y * y)) - (x * y)) * z

```

And another excellent solution from Jacob Mikkelsen (also minimising instructions and registers):


```
# Rewrite as (x(x - y) + y^2) * z
lw  $t0, x          # x
lw  $t1, y          # y
sub $v0, $t0, $t1    # x - y
mul $v0, $v0, $t0    # x(x - y) = x^2 - xy
mul $t1, $t1, $t1    # y^2
add $v0, $v0, $t1    # x^2 + y^2 - xy
lw  $t0, z          # z
mul $v0, $v0, $t0    # (x^2 + y^2 - xy) * z
```

COMP1521 20T2: Computer Systems Fundamentals is brought to you by
the [School of Computer Science and Engineering](#)
at the [University of New South Wales](#), Sydney.

For all enquiries, please email the class account at cs1521@cse.unsw.edu.au

CRICOS Provider 00098G