

Week 10 Tutorial Questions

1. How is the assignment going?

Does anyone have hints or advice for other students?

Has anyone discovered interesting cases that have to be handled?

2. Write a C program, `print_diary.c`, which prints the contents of the file `$HOME/.diary` to stdout

The lecture example [getenv.c](#) shows how to get the value of an environment variable.

`snprintf` is a convenient function for constructing the pathname of the diary file.

3. Write a C program, `print_file_bits.c`, which given as a command line arguments the name of a file contain 32-bit hexadecimal numbers, one per line, prints the low (least significant) bytes of each number as a signed decimal number (-128..127).

4. Assume we have 6 virtual memory pages and 4 physical memory pages and are using a least-recently-used (LRU) replacement strategy.

What will happen if these virtualmemory pages were accessed?

5 3 5 3 0 1 2 2 3 5

5. Discuss code supplied for the *lru* lab exercise.

```

// Simulate LRU replacement of page frames

#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <assert.h>

// represent an entry in a simple inverted page table

typedef struct ipt_entry {
    int virtual_page;        // == -1 if physical page free
    int last_access_time;
} ipt_entry_t;

void lru(int n_physical_pages, int n_virtual_pages);
void access_page(int virtual_page, int access_time, int n_physical_pages, struct ipt_entry *ipt);

int main(int argc, char *argv[]) {
    if (argc != 3) {
        fprintf(stderr, "Usage: %s <n-physical-pages> <n-virtual-pages>\n", argv[0]);
        return 1;
    }
    lru(atoi(argv[1]), atoi(argv[2]));
    return 0;
}

void lru(int n_physical_pages, int n_virtual_pages) {
    printf("Simulating %d pages of physical memory, %d pages of virtual memory\n",
           n_physical_pages, n_virtual_pages);
    struct ipt_entry *ipt = malloc(n_physical_pages * sizeof *ipt);
    assert(ipt);

    for (int i = 0; i < n_physical_pages; i++) {
        ipt[i].virtual_page = -1;
        ipt[i].last_access_time = -1;
    }

    int virtual_page;
    for (int access_time = 0; scanf("%d", &virtual_page) == 1; access_time++) {
        assert(virtual_page >= 0 && virtual_page < n_virtual_pages);
        access_page(virtual_page, access_time, n_physical_pages, ipt);
    }
}

// if virtual_page is not in ipt, the first free page is used
// if there is no free page, the least-recently-used page is evicted
//
// a single line of output describing the page access is always printed
// the last_access_time in ipt is always updated

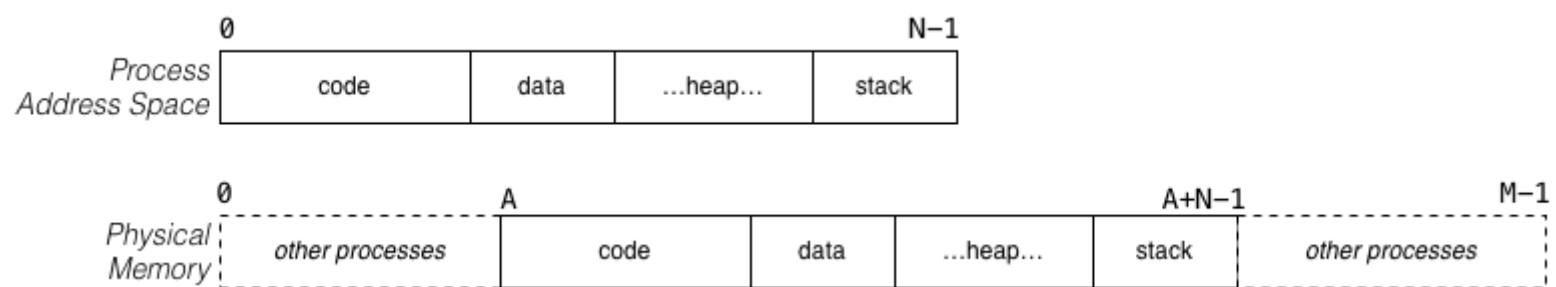
void access_page(int virtual_page, int access_time, int n_physical_pages, struct ipt_entry *ipt) {

    // PUT YOUR CODE HERE TO HANDLE THE 3 cases
    //
    // 1) The virtual page is already in a physical page
    //
    // 2) The virtual page is not in a physical page,
    //    and there is free physical page
    //
    // 3) The virtual page is not in a physical page,
    //    and there is no free physical page
    //
    // don't forgot to update the last_access_time of the virtual_page

    printf("Time %d: virtual page %d accessed\n", access_time, virtual_page);
}

```

6. Each new process in a computer system will have a new address space. Which parts of the address space contain initial values at the point when the process starts running? Code? Data? Heap? Stack? Which parts of the address space can be modified as the process executes?
7. One possible (and quite old) approach to loading programs into memory is to load the entire program address space into a single contiguous chunk of RAM, but not necessarily at location 0. For example:

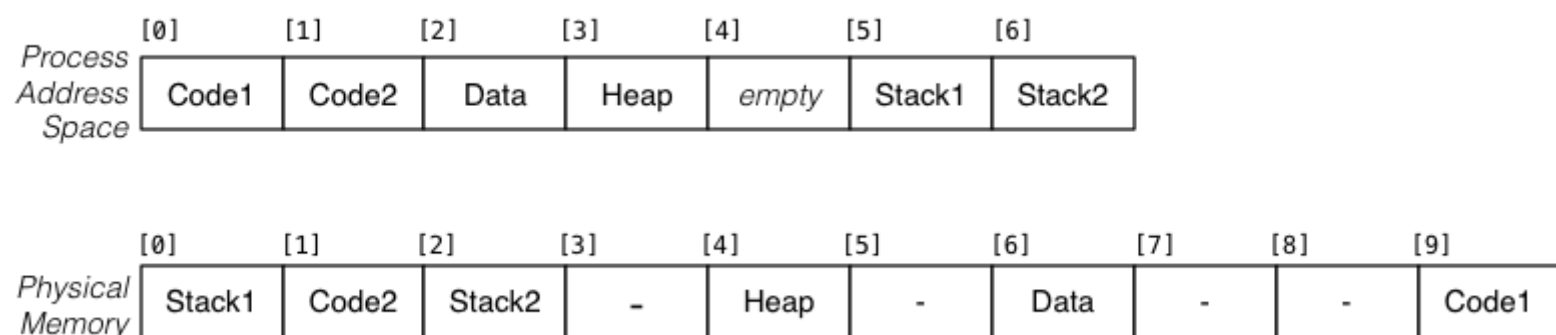


Doing this requires all of the addresses in the program to be rewritten relative to the new base address.

Consider the following piece of MIPS code, where `loop1` is located at `0x1000`, `end_loop1` is located at `0x1028`, and `array` is located at `0x2000`. If the program containing this code is loaded starting at address `A = 0x8000`, which instructions need to be rewritten, and what addresses are in the relocated code?

```
li $t0, 0
li $t1, 0
li $t2, 20 # elements in array
loop1:
bge $t1, $t2, end_loop1
mul $t3, $t1, 4
la $t4, array
add $t3, $t3, $t4
lw $t3, ($t3)
add $t1, $t1, $t3
add $t1, $t1, 1
j loop1
end_loop1:
```

8. What is the difference between a *virtual address* and a *physical address*?
9. Consider a process whose address space is partitioned into 4KB pages and the pages are distributed across the memory as shown in the diagram below:



The low byte address in the process is 0 (in Code1) and the top byte address in the process is 28671 (max address in page containing Stack2).

For each of the following process addresses (in decimal notation), determine what physical address it maps to.

- `jal func`, where the label `func` is at 5096
- `lw $s0, ($sp)`, where `$sp` contains 28668
- `la $t0, msg`, where the label `msg` is at 10192

10. The *working set* of a process could be defined as the set of pages being referenced by the process over a small window of time. This would naturally include the pages containing the code being executed, and the pages holding the data being accessed by this code.

Consider the following code, which computes the sum of all values in a very large array:

```
int bigArray[100000];
// ...
int sum = 0;
for (int i = 0; i < 100000; i++)
    sum += bigArray[i];
```

Answer the questions below under the assumptions that pages are 4 KiB (4096 bytes), all of the above code fits in a single page, the `sum` and `i` variables are implemented in registers, and there is just one process running in the system.

- How large is the working set of this piece of code?

b. Assuming that the code is already loaded in memory, but that none of `bigArray` is loaded, and that *only* the working set is held in memory, how many page faults are likely to be generated during the execution of this code?

11. Consider a (very small) virtual memory system with the following properties:

- a process with 5 pages
- a memory with 4 frames
- page table entries containing (Status, MemoryFrameNo, LastAccessTime)
- pages status is one of NotLoaded, Loaded, Modified (where Modified implies Loaded)

Page table:

Page Table

[0]	[1]	[2]	[3]	[4]
Status	Status	Status	Status	Status
FrameNo	FrameNo	FrameNo	FrameNo	FrameNo
LastAccess	LastAccess	LastAccess	LastAccess	LastAccess

indexes
are page
numbers

If all of the memory frames are initially empty, and the page table entries are flagged as NotLoaded, show how the page table for this process changes as the following operations occur:

```
read page0,  read page4,  read page0,  write page4,  read page1,
read page3,  read page2,  write page2,  read page1,  read page0,
```

Assume that a LRU page replacement policy is used, and unmodified pages are considered for replacement before modified pages. Assume also that access times are clock ticks, and each of the above operations takes one clock tick.

12. Some commands on Unix allow you to name the files that they operate on, e.g.,

```
$ cat file
```

Commands that read from their standard input allow you to specify which file they read their input from by redirecting their standard input, e.g.,

```
$ cat < file
```

Describe how each of these cases might be implemented. Assume that once the file is made accessible, it is scanned and copied to the standard output (file descriptor 1, or the `#define'd` constant `STDOUT_FILENO`) as follows:

```
int infd;           // input file descriptor
char buffer[BUFSIZ]; // input file buffer
int nread;          // # characters read

while ((nread = read (infd, buffer, BUFSIZ)) > 0)
    write (STDOUT_FILENO, buffer, nread);
```

Revision questions

The following questions are primarily intended for revision, either this week or later in session. Your tutor may still choose to cover some of these questions, time permitting.

19. Write a C program, `compile.c`, which is given 1+ command-line arguments which are the pathname of single file C programs. It should compile each program with `dcc`.

It also should print the compile command to `stdout`.

```
$ dcc compile.c -o compile
$ ./compile file_sizes.c file_modes.c
/usr/local/bin/dcc file_modes.c -o file_modes
/usr/local/bin/dcc file_sizes.c -o file_sizes
```

Make sure you handle errors, for example, you should stop if any compile fails.

20. Consider the following edited output from the [ps\(1\)](#) command running on one of the CSE servers:

PID	VSZ	RSS	TTY	STAT	START	TIME	COMMAND
1	3316	1848	?	Ss	Jul08	1:36	init
321	6580	3256	pts/52	Ss+	Aug26	0:00	-bash
334	41668	11384	pts/44	Sl+	Aug02	0:00	vim timing_result.txt
835	6584	3252	pts/124	Ss+	Aug27	0:00	-bash
857	41120	10740	pts/7	Sl+	Aug22	0:00	vi echon.pl
924	6524	3188	pts/184	Ss	15:52	0:00	-bash
938	3664	96	pts/184	S	15:52	0:00	/usr/local/bin/checkmail
1199	6400	3004	pts/142	Ss	Oct05	0:00	-bash
1381	41504	11436	pts/142	Sl+	Oct05	0:00	vim PageTable.h
2558	3664	96	pts/120	S	13:47	0:00	/usr/local/bin/checkmail
2912	41512	11260	pts/46	Sl+	Aug02	0:00	vim IntList.c
3483	14880	5168	pts/149	S+	Sep20	0:00	gnuplot Window.plot
3693	41208	11240	pts/120	Tl	13:50	0:00	vim trace4
3742	6580	3320	pts/116	Ss+	Sep07	0:00	-bash
5531	6092	2068	pts/158	R+	16:04	0:00	ps au
5532	4624	684	pts/158	S+	16:04	0:00	cut -c10-15,26-
5538	3664	92	pts/137	S	15:05	0:00	/usr/local/bin/checkmail
6620	5696	3028	pts/89	S+	Aug13	0:00	nano PingClient.java
7132	41516	11196	pts/132	Sl+	Sep08	0:00	vim board1.s
12256	335316	10436	?	Sl	Aug14	15:01	java PingServer 3331
12272	4260	2816	?	Ss	Aug02	10:34	tmux
12323	10276	4564	?	S	Sep09	0:02	/usr/lib/i386-linux-gnu/gconf/gconfd-2
12461	4260	2808	?	Ss	Sep02	5:42	tmux
13051	43448	13320	pts/110	Sl+	Sep05	0:02	vim frequency.pl
13200	47772	21928	?	Ssl	15:19	0:02	gvim browser.cgi
13203	41756	11560	pts/26	Sl+	Aug12	0:02	vim DLList.h
13936	11872	6856	?	S	Sep19	0:06	/usr/lib/gvfs/gvfs-gdu-volume-monitor
30383	7624	3828	pts/77	S+	Aug23	336:28	top

- Where might you look to find out the answers to the following questions?
- What does each of the columns represent?
- What do the first characters in the STAT column mean?
- Which process has consumed the most CPU time?
- Why do some processes have no TTY?
- When was this machine last re-booted?

21. The Unix/Linux shell is a text-oriented program that runs other programs. It behaves more-or-less as follows:

```

print a prompt
while (read another command line) {
    break the command line into an array of words (args[])
    // args[0] is the name of the command, a[1],... are the command-line args
    if (args[0] starts with '.' or '/')
        check whether args[0] is executable
    else
        search the command PATH for an executable file called args[0]
    if (no executable called args[0])
        print "Command not found"
    else
        execute the command
    print a prompt
}

```

- How can you find what directories are in the PATH?
- Describe the “search the command PATH” process in more detail. What the kinds of system calls would be needed to determine whether there was an executable file in one of the path directories?

22. The [kill\(1\)](#) command (run from the shell command-line) and the `kill()` system call can be used to send any of the defined signals to a specified process. For each of the following signals, explain the circumstances under which it might be generated (apart from `kill(1)`), and what is the default effect on the process receiving the signal:

- SIGHUP
- SIGINT
- SIGQUIT
- SIGABRT
- SIGFPE

f. SIGSEGV

i. SIGSEGV

g. SIGPIPE

h. SIGTSTP

i. SIGCONT

23. The *sigaction(2)* function for defining signal handlers takes three arguments:

- `int signum` ... the signal whose handler is being defined
- `struct sigaction *act` ... pointer to a record describing how to handle the signal
- `struct sigaction *oldact` ... pointer to a record describing how the signal was handled (set by *sigaction(3)* if not NULL)

The `struct sigaction` record includes a field of type `void (*sa_handler)(int)`.

Describe precisely what this field is, and what its type signature means.

24. Consider the following program:

```
// assume a bunch of #include's

static void handler (int sig)
{
    printf ("Quitting...\n");
    exit (0);
}

int main (int argc, char *argv[])
{
    struct sigaction act;
    memset (&act, 0, sizeof (act));
    act.sa_handler = &handler;
    sigaction (SIGHUP, &act, NULL);
    sigaction (SIGINT, &act, NULL);
    sigaction (SIGKILL, &act, NULL);
    while (1)
        sleep (5);
    return 0;
}
```

What does this program do if it receives

- a. a SIGHUP signal?
- b. a SIGINT signal?
- c. a SIGTSTP signal?
- d. a SIGKILL signal?

COMP1521 20T2: Computer Systems Fundamentals is brought to you by
the [School of Computer Science and Engineering](#)
at the [University of New South Wales](#), Sydney.
For all enquiries, please email the class account at cs1521@cse.unsw.edu.au

CRICOS Provider 00098G